

CONCURRENT PERMISSION MACHINE FOR
MODULAR PROOFS OF OPTIMIZING COMPILERS
WITH SHARED MEMORY CONCURRENCY.

SANTIAGO CUELLAR

A DISSERTATION
PRESENTED TO THE FACULTY
OF PRINCETON UNIVERSITY
IN CANDIDACY FOR THE DEGREE
OF DOCTOR OF PHILOSOPHY

RECOMMENDED FOR ACCEPTANCE
BY THE DEPARTMENT OF
COMPUTER SCIENCE
ADVISER: ANDREW APPEL

SUBM. DATE

© Copyright by Santiago Cuellar, tbd copyright.

All rights reserved.

Abstract

Optimizing compilers change a program based on a formal analysis of its code and modern processors further rearrange the program order. It is hard to reason about such transformations, which makes them a source of bugs, particularly for concurrent shared-memory programs where the order of execution is critical. On the other hand, programmers should reason about their program in the source language, which abstracts such low level details.

We present the Concurrent Permission Machine (CPM), a semantic model for shared-memory concurrent programs, which is: (1) sound for high-order Concurrent Separation Logic, (2) convenient to reason about compiler correctness, and (3) useful for proving reduction theorems on weak memory models. The key feature of the CPM is that it exploits the fact that correct shared-memory programs are permission coherent: threads have (at any given time) noncompeting permission to access memory, and their load/store operations respect those permissions.

Compilers are often written with sequential code in mind, and proving the correctness of those compilers is hard enough without concurrency. Indeed, the machine-checked proof of correctness for the CompCert C compiler was a major advance in the field. Using the CPM to conveniently distinguish sequential execution from concurrent interactions, I show how to reuse the (sequential) CompCert proof, without major changes, to guarantee a stronger concurrent-permission-aware notion of correctness.

Acknowledgements

Acknowledgements...

Dedication.

Contents

Abstract	iii
Acknowledgements	iv
List of Tables	viii
List of Figures	x
1 Introduction	1
2 Related Work	2
2.1 Past work 1	2
3 Top to Bottom structure	3
3.1 Main theorem	3
3.2 The concurrent permission machine	4
3.3 Limitations of the CPM model	4
4 Making Compcert a <i>conjunctive</i> compiler	5
4.1 Passing arguments to main.	9
4.1.1 Simulating initialization preprocessing.	10
4.2 Memory events.	13
4.3 Conjunctive Semantics	16
4.4 More revealing simulations	17

5	Compiler Correctness	24
5.1	Compiler Theorem	24
5.1.1	Self simulations	24
5.2	Compiler Theorem Implementation	24
6	Conclusion	25
6.1	Future Work	25
A	Implementation Details	26
A.1	CPM design	26
	Bibliography	27

List of Tables

4.1	The function <code>remember</code> records the address of some buffer, and <code>incr</code> increments it by one.	5
4.2	The <code>initial_state</code> in C and Clight is a call to <code>main</code> . It also enforces that it takes no arguments (<code>Tnil</code>) and returns an integer (<code>type_int32s</code>). . . .	6
4.3	The events in CompCert	7
4.4	Percentage change to CompCert: changes are calculated from the number of lines added as given by running <code>git diff</code> between our version of CompCert and the master branch. For each feature, an estimated percent is provided.	8
4.5	Entry simulation diagrams. (a) if s_1 is an initial state for the source program, then there exists some state s_2 that is related to s_1 and is an initial state for the compiled program. (b) Just like the diagram for initial states, but it generalizes and exposes the initial memory m_0 , the entry function f and the arguments <code>args</code> . The entire simulation is parametric on the realation \sim	10
4.6	The <code>entry_point</code> predicate in Clight	12
4.7	Part of the <code>entry_point</code> predicate in Mach	13
4.8	The new events in CompCert: <code>mem_effect</code> reflects changes to memory and <code>delta_perm_map</code> represents trasfer of Cur permissions.	14

4.9	<code>at_external</code> definition for Clight. The function checks that (1) the current state is about to make a function call, (2) that the function is an External function and, (3) that the external function cannot be inlined (The compiler is allowed to inline specific functions such as <code>memcpy</code> and certain builtins).	16
4.10	Step simulation step diagrams. (a) if s_1 takes a step to s_2 with trace t and s_1 is related to some s_2 , then there s_2 can take a number of steps with trace t to a new state s'_2 related to s'_1 . (b) The new diagram exposes the memory reordering injections j and j' and the traces t and t' are equivalent up to injection, by <code>inject_trace.strong j' t t'</code> .	18
4.11	At external step diagram (<code>simulation_atx</code>). Exclusive for external function calls, this diagram follows the simulation diagram in 4.10, but enforces that the compiled execution takes only one step.	20
4.12	Stronger simulation for external steps, that universally quantifies over all injected traces. Lines 8-11 describe the existentially quantified diagram as described in 4.11. Lines 12-15, in bold, describe all the other executions that may have undefined values determined. <code>inject_trace</code> is the predicate that allows undefined values to be mapped to defined ones.	21

List of Figures

Chapter 1

Introduction

The main ultimate goal of this thesis is to create a modular way to prove compilers safe, even in the presence of concurrency. There are already several compilers proven safe, with machine checked proofs, that give guarantees for sequential programs. Instead of reinventing the wheel, I intend to leverage those proofs and use it to give guarantees for concurrent programs.

On top of "just a compiler proof", our system is intended to be compatible with a realistic separation logic and usable on realistic machines like X86 TSO.

The main result:

Contributions: In this thesis I have....

Chapter 2

Related Work

Here I talk about related work:

2.1 Past work 1

Topic 1.

Chapter 3

Top to Bottom structure

Note.

3.1 Main theorem

Explaining the main theorem:

Assumptions:

Theorem `top2bottom_correct`:

`CSL_correct C_program →`

`CompCert_compiler C_program = Some Asm_program →`

$\forall (\text{src_m}:\text{Memory.mem}) (\text{src_cpm}:\text{Clight.state}),$

$\text{CSL_init_setup } C_program \text{ src_m src_cpm} \rightarrow$

$\text{Clight.entry_point } (\text{Clight.globalenv } C_program) \text{ src_m src_cpm } (\text{main_ptr } C_program) \text{ nil} \rightarrow$

$\forall (\text{limited_builtins}:\text{Asm_core.safe_genv } \text{x86_context.X86Context.the_ge}),$

$\text{asm_prog_well_formed } \text{Asm_program } \text{limited_builtins} \rightarrow$

$\forall (\text{U}:\text{schedule}), \exists (\text{tgt_m0 } \text{tgt_m}:\text{mem}) (\text{tgt_cpm}:\text{ThreadPool.t}),$

$\text{permissionless_init_machine}$

$\text{Asm_program } \text{limited_builtins}$

$\text{tgt_m0 } \text{tgt_cpm } \text{tgt_m } (\text{main_ptr } C_program) \text{ nil} \wedge$

$\text{spinlock_safe } \text{U } \text{tgt_cpm } \text{tgt_m}.$

- `CSL_correct C_program`: States that the program has been proven correct in CSL, as described in [section CSL ref]
- `CompCert_compiler`: States that the CompCert compiler translates `C_program` into the assembly program `Asm_program`.
- `CSL_init_setup`: States that `src_m` and `src_cpm` are the initial memory and state as defined by the program `C_program`... According to CSL
- `Clight.entry_point` : Same as above... According to Clight.
- `limited_builtins`: Statically checkable property stating that `Asm_program` doesn't have unsupported builtins.
- `asm_prog_well_formed`: Another statically checkable property stating that the initial memory and global environment created by `Asm_program` are well formed.

Conclusions:

- `permissionless_init_machine`: There exists initial memories and state for `Asm_program`
- `spinlock_safe`: The initial state is safe and "spinlock well synchronized" as defined in [reference to spinlock]

3.2 The concurrent permission machine

Describe CPM.

3.3 Limitations of the CPM model

What are the limitations? - That coin example from Ernie Cohen. - Angel transferring too many permissions proof.

Chapter 4

Making CompCert a *conjunctive* compiler

The simple code in 4.1 communicates with its environment in two main ways: (1) it takes an address as input and (2) reads from and writes to this location to increment the value stored there. We will see how the specifications of CompCert prevents us from reasoning about such program as compilation units and as external functions and we will show how to extend the specificatoin of a compiler to lift thees limitations.

First, CompCert can't give any guarantees about compiling the code in 4.1 because it is not a complete program. It is reasonable to expect that the function **remember** runs safely, given some assumptions (e.g. ***p** is a valid address in memory). Unfortu-

```
1 int *buff;
2 void remember(int *p){
3     buff = p;
4 }
5 void incr(void){
6     ++(* buff);
7 }
```

Table 4.1: The function **remember** records the address of some buffer, and **incr** increments it by one.

Inductive `initial_state` (`p`: program): state $\rightarrow \mathbb{P} :=$
| `initial_state_intro`: $\forall b\ f\ m0,$
 let `ge` := `Genv.globalenv p` **in**
 `Genv.init_mem p` = Some `m0` \rightarrow
 `Genv.find_symbol ge p.(prog_main)` = Some `b` \rightarrow
 `Genv.find_funct_ptr ge b` = Some `f` \rightarrow
 `type_of_fundef f` = `Tfunction Tnil type_int32s cc_default` \rightarrow
 `initial_state p` (`Callstate f nil Kstop m0`).

Table 4.2: The `initial_state` in C and Clight is a call to `main`. It also enforces that it takes no arguments (`Tnil`) and returns an integer (`type_int32s`).

nately, CompCert’s semantics assumes that a program starts executing with a call to `main()` with no arguments. In fact, the only possible initial states, are characterized by a predicate `initial_state: state $\rightarrow \mathbb{P}$` that takes no additional arguments. You can see an instantiation of the predicate for Clight in 4.2. So, even though CompCert correctly compiles the code, it’s specification gives no guarantees of any execution other than the one that starts by calling `main` with no arguments.

Second, CompCert’s can’t reason about programs that call `remember` and `incr` because their behavior depends on their internal state (the `buff` pointer). CompCert’s semantics allows calls to external calls that are assumed to be correct but unfortunately the specification of ”correctness” is too strict, it assumes that the function’s behavior is fully determined by (1) the state of memory, (2) the function arguments and (3) the events produced by the function.¹ The execution of `incr` also depends on arguments passed to a previous external call (namely, `*p` passed to `remember`). One could make this explicit in the event trace but CompCert events, shown in 4.3, can only contain integers, floats or pointers to global variables.

Moreover, in the CompCert semantics, the entire behavior of external functions is bundled into one big step. From looking at an execution internal steps and external

¹Leroy [2] claims that ”inputs given to the programs are uniquely determined by their previous outputs”, but this is not exactly correct. A more accurate representation would be to say ”inputs given to the programs are uniquely determined by their last outputs”. As we will see in 4.4, it would be much stronger to determine inputs based on all historic outputs.

Inductive event: Type :=	Inductive eventval: Type :=
Event_syscall: string → list eventval → eventval → event	EVint: int → eventval
Event_vload: memory_chunk → ident → ptrofs → eventval → event	EVlong: int64 → eventval
Event_vstore: memory_chunk → ident → ptrofs → eventval → event	EVfloat: float → eventval
Event_annot: string → list eventval → event	EVsingle: float32 → eventval
	EVptr_global: ident → ptrofs → eventval.

Table 4.3: The events in CompCert

function calls are uniform. This consistency is very useful when reasoning about the compilation of the program, where we want to abstract external calls. Nevertheless, when reasoning about a program in a context, it is more useful to replace the big step external calls, with their small step semantics. Regrettably, the specification of CompCert does not even guarantee that the source and target programs call the same external functions. In theory, CompCert could replace an external function call with internal steps as long as they had the same (possibly empty) trace. In practice, obviously, CompCert does not do that, but it is not exposed in its specification.

Finally, the correctness of CompCert is stated as semantic preservation theorem, where the traces are the preserved behavior. The proof of this semantic preservation, uses a forward simulation². We find it more convenient to expose this simulation as the specification of the compiler, deriving semantic preservation as a corollary. This way we can easily express the relation between external function calls in source and target, as simulation diagrams. Moreover, we can expose the relation between memories, in source and target.

We move, then, to lift this limitations according to the following richer notion of specification

²Forward simulation and determinism of the target language implies bisimulation and thus preservation of behavior.

	Percent change
Arguments in <code>main</code>	0
Injectable Traces	0
Semantics	0
Simulations	0
Total	x

Table 4.4: Percentage change to CompCert: changes are calculated from the number of lines added as given by running `git diff` between our version of CompCert and the master branch. For each feature, an estimated percent is provided.

Definition 1 (Conjunctive compiler specification) *We say that a compiler’s specification is conjunctive if it satisfies the following*

- *The execution of a program can start in any of its public functions and all the public functions, including `main`, can take arguments.*
- *The execution trace supports events that can describe changes to memory. Since memory may be rearranged through compilation, we call these memory events.*
- *There are functions `at_external` and `after_external` that describe states before and after external calls. `at_external` extracts the external function being called and its arguments. `after_external` constructs the state after the external function executes. Also, source and target languages operate over the same memory and the memory can be separated by `get_mem` function. We call these exposed semantics.*
- *The correctness of the compiler is stated as a forward simulation and the semantic preservation derive as a corollary. The forward simulations preserve external function calls and exposes the relation between source and target memories, throughout the execution. We call this an exposed simulation.*

In the rest of the chapter, we describe how we have improved the specification of CompCert to make it conjunctive. We first describe how to generalize `initial_state` to accept functions with arguments and other than `main`. Second we describe how

to add memory events to CompCert. Then we show how to easily define exposed semantics for every language in CompCert and finally we show how to define and prove exposed simulations.

The changes described here represent only a x% change to CompCert, as measured by running `git diff` in CompCert before and after our changes. The amount changed for every feature proposed is described in Table 4.4.

4.1 Passing arguments to main.

CompCert can compile programs where `main` takes arguments, but its correctness theorem gives no guarantees about their translation. That is because its semantic model assumes that `main` takes no arguments (See 4.2), but real C programs can take up to two arguments `argc` (argument count) and `argv` (argument vector). Also, all execution in the semantics of CompCert start with a call to `main()`, but `main` is nothing but an agreed upon term for startup code. Furthermore, the semantics assumes that the initial memory contains only the global environments. In this section we define a new predicate `entry_point`, generalizing `initial_state`, which accepts any public function, including those with arguments. The predicate also admits other memories containing, for example, the stack of other processes.

Passing arguments to main is of general interest and particularly important for our work with concurrency. Spawning new threads behaves very similar to starting a program by calling `main`: The library function that spawns a new thread `f` (e.g. `pthread_create`) must create a new stack, push the arguments to stack and then call `f` just like a program startup (which calls `_start()` and `__libc_start_main()`). Our new predicate `entry_point` is general enough to capture this two types of preprocessing. Notice that if we restricted our semantics to spawning threads with no arguments, threads wouldn't be able to share pointers and thus they would all execute in disjoint

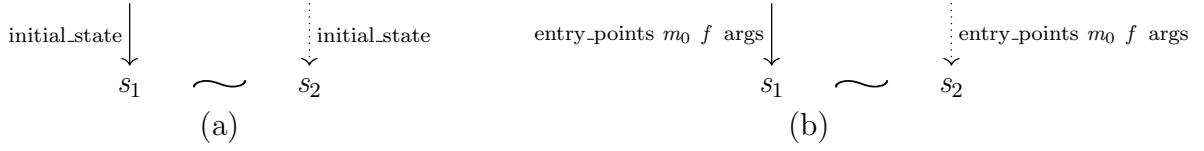


Table 4.5: Entry simulation diagrams. (a) if s_1 is an initial state for the source program, then there exists some state s_2 that is related to s_1 and is an initial state for the compiled program. (b) Just like the diagram for initial states, but it generalizes and exposes the initial memory m_0 , the entry function f and the arguments args . The entire simulation is parametric on the relation \sim .

pieces of memory with no communication. That would be a much easier and less interesting result.

Now that we defined a new starting point for executions, we must prove that compilation preserves the predicate `entry_point` in the same way that CompCert’s simulation preserves `initial_state`. The proof largely follows the simulation of internal function calls which is already proven in CompCert, so we omit the details here.

4.1.1 Simulating initialization preprocessing.

Before `main` ever starts executing, there is an initialization process that sets up the stack and the registers, before `main` is executed. The preprocessing function(s) often called `_start`, `premain` or `startup`, will create a stack, set up the arguments in the stack or in registers, set up the environment (`envp`), set up the return address (a call to `exit()`), and other bookkeeping.

One of the difficulty in simulating the initialization is that, in architectures such as *x86* in 32-bit mode, all arguments are passed in memory. As the comments in the CompCert code would put it ”Snif!” [3]. Even architectures that allow argument passing in registers, such as *x86* in 64-bit mode, have a limited number of registers and will pass arguments on the stack after those run out. This ”pre-stack”, that contains the arguments, does not correspond to any function in the call stack of the

program; it corresponds to the stack frame of `_start()`, which is part of the linked program.

Our generic predicate `entry_points: mem \rightarrow state \rightarrow val \rightarrow list val \rightarrow \mathbb{P}` takes a memory `m`, a state `s`, a pointer to the entry function `fun_ptr` of type `val`, and a list of arguments `args`. This predicate is language dependent, but it generally has three parts:

1. Checks memory is well formed. That is, it contains no ill-formed pointers to invalid memory. CompCert generally maintains that well formed programs don't create dangling pointers pointers.³
2. Arguments are well formed. Among other things, they have the right types for the function being called, they have no ill-formed pointers, and fit in the stack.
3. Global environment is allocated correctly. This makes sure that the environment is in memory and that the function is declared as a public in the environment.

In the rest of this section, we explore the definition of `entry_points` for different languages and show how we prove that different CompCert phases preserve the predicate.

C frontend

All of the C-like languages (*Clight*, *Csharp*, *Csharpminor*) have similar `entry_point`, so we present here the one for *Clight* in 4.6. Lines 4-6 ensure that the environment is allocated in memory and it contains the function `f` with the right type signature. Line 7 states that the initial memory has no dangling pointers. Lines 8-11 say that the arguments have the right type and have no dangling pointers. The predicate `bounded-args`, enforces that the arguments fit in the stack, which is architecture dependent. It is reasonable to replace this with a small enough bound that fits all

³A well formed program, should not compare, read or write to invalid pointers. Hence, dangling pointers behave semantically as undefined values and could be modeled that way.

```

1 Inductive entry_point (ge:genv): mem → state → val → list val →  $\mathbb{P}$  :=
2 | initi_core:  $\forall f \text{ fb } m0 \text{ args } \text{targs},$ 
3   let sg:= signature_of_type targs type_int32s cc_default in
4   type_of_fundef (Internal f) = Tfunction targs type_int32s cc_default →
5   Genv.find_funct_ptr ge fb = Some (Internal f) →
6   globals_not_fresh ge m0 →
7   Mem.mem_wd m0 →
8   Val.has_type_list args (tylist_of_typelist targs) →
9   vars_have_type (fn_vars f) targs →
10  vals_have_type args targs →
11  Mem.arg_well_formed args m0 →
12  bounded_args sg →
13  entry_point ge m0 (Callstate (Internal f) args (Kstop targs) m0).

```

Table 4.6: The `entry_point` predicate in Clight

architectures, like 4, but we keep it general here. Finally, the entry state, in line 13, is defined as a call to `f` with an empty continuation.

Our empty continuation `Kstop` takes `targs` as an argument. That is because continuations also represent the call stack; `Kstop targs` represents the "pre-stack" of `_start()` that might contain some arguments for `f`.

Register transfer languages

In the Cminorgen phase, CompCert coalesces all function variables into a stackframe. Some functions might get empty stackframes (i.e., a memory block with empty permissions, that cannot be written to), if none of their variables has their address taken. These stackframes are important, even the empty ones, because that is where spill variables will be written after register allocation in the Allocation phase. We follow suit and create an empty stackframe for `_start()`. The stack is empty because the compiler has not yet decided what arguments will be passed in memory and which ones in register. Even for architectures that pass all arguments in memory, this is not until the Stacking pass. So for languages before Stacking (Cminor, Cmi-

```

1      :
2      let '(stk_sz,ret_ofs,parent_ofs) := stack_defs (fn_sig f) in
3      Mem.alloc m0 0 stk_sz = (m1, spb) →
4      let sp:= Vptr spb Ptrofs.zero in
5      store_stack m1 sp Tptr parent_ofs Vnullptr = Some m2 →
6      store_stack m2 sp Tptr ret_ofs Vnullptr = Some m3 →
7      make_arguments (Regmap.init Vundef) m3 sp
8              (loc_arguments (funsig (Internal f))) args = Some (rs, m4) →
9      :

```

Table 4.7: Part of the `entry_point` predicate in Mach

norSel, RTL, LTL, Linear), there is an extra line in `entry_point` to make sure that the empty stackframe is allocated:

Mem.alloc m0 0 0 = (m1, stk)

The rest of the predicate is almost identical to the one in 4.6.

Machine languages

In the machine languages (Mach and Asm), a function expects certain shape from the stackframe of its caller. We replace the empty stackframe allocation above, with the constructon of the "pre-stackframe" as shown in 4.7. The function `stack-defs` is an architecture dependent function that calculates the layout of the "pre-stack" and returns the size `stk_sz`, the ofset of the return address `ret_ofs` and a back link to parent frame `parent_ofs`. The last two values are unused, but the stack must have space for them. Line 3 allocates the stack of the correct size. Lines 5-8 store the return address, the link to the parent and the arguments in the stack.

4.2 Memory events.

The visible behavior of the CompCert semantics (for all languages) is a trace of events, as described in 4.3. It records interactions with the outside world; for example, the

Inductive event: Type :=	Inductive mem_effect: Type :=
...	Write : $\forall (b : \text{block}) (ofs : Z)$
Event_acq_rel:	(bytes : list memval), mem_effect
list mem_effect \rightarrow	Alloc: $\forall (b : \text{block})(lo\ hi:Z), \text{mem_effect}$
delta_perm_map \rightarrow	Free: $\forall (l: \text{list } (\text{block} * Z * Z)), \text{mem_effect}.$
list mem_effect $\rightarrow \text{event}$	
Event_spawn:	
block \rightarrow	
delta_perm_map \rightarrow	
delta_perm_map $\rightarrow \text{event}.$	

Table 4.8: The new events in CompCert: **mem_effect** reflects changes to memory and **delta_perm_map** represents transfer of **Cur** permissions.

results of a read system call will record a **Event_syscall** together with the name of the system call, its parameters, and its result. The only pointers that can appear in the trace are locations of global variable. As we explained above, these events are insufficient to capture the behavior of **incr** in 4.1. The events are also not sufficient to express any kind of behavior that depends on the memory. For example, a call to **pthread_mutex_unlock(&l)** not only changes the state of the lock **l**, conceptually it gives away control to the data in memory protected by **l**. Such behavior that reveals locations in memory cannot be expressed in CompCert events that can't have pointers on them (except the location of global variables). More generally, these events are poorly suited to express any sort of shared memory interactions, such as concurrency or separated compilation. We propose to include 2 new types of events which we call *memory events* as described in table 4.8. As their name suggests, they contain references to locations in memory, beyond the global variables. The first one, **Event_acq_rel**, represents a generic memory interactions where the external function performs some arbitrary changes to memory, recorded by a list of **mem_effects**, and transfers some permissions recorded by **delta_perm_map**. We find it convenient to split the effects on memory as those that happen before the changes in permissions and those that happen after. The second one, **Event_spawn**, represents the creation of a new thread or a new module. It records the function being called, as a block

number, and the change in permissions by two `delta_perm_map`, one representing the permissions given and the other the starting permissions of a new thread/module.

The correctness of the CompCert compiler is formulated as a preservation of traces. However, if our traces contain memory events and the memory can be reordered by compilation, the trace must be reordered accordingly. Thus our new correctness will be formulated up to memory reordering as shown in ???. In fact, if a compiler pass (or passes) don't change the order of memory (such as equality and extension passes), then the trace is preserved and nothing changes in the simulation. If the pass changes the order of memory according to some j , then we need show that the source and target traces, t and t' , are related by `inject_trace_strong j t t'` (denoted $t \xrightarrow{j'} t'$). That means that events in t and t' are identical except all memory locations are reordered according to j . To add this property, we need simulations to expose how they reorder memory. The full description of our new conjunctive simulation will be explained in section 4.4.

Compilers not only reorder memory, sometimes undefined values in the source program are mapped to concrete values in the target. Such mappings are useful in compiler passes such as register coalescing, where registers that were previously uninitialized, can now map to an initialized register with concrete values. The predicate `inject_trace_strong` always maps undefined values to undefined values. This is a reasonable restriction since inspecting undefined values is not an allowed behavior so they shouldn't appear in the trace. However, if an application required traces with undefined values, we can still support that. It turns out that it is enough to prove that the execution with the *strongly injected* trace is safe, and from it we can derive safe executions for all traces that have some undefined values defined.

Definition `at_external (c: state) : option (external_function * list val) :=`
`match c with`
`| Callstate fd args k _ =>`
`match fd with`
`| External ef targ tres cc => if ef.inline ef then None else Some (ef, args)`
`| _ => None`
`end`
`| _ => None`
`end.`

Table 4.9: `at_external` definition for Clight. The function checks that (1) the current state is about to make a function call, (2) that the function is an External function and, (3) that the external function cannot be inlined (The compiler is allowed to inline specific functions such as `memcpy` and certain builtins).

4.3 Conjunctive Semantics

As described in the introduction, we need more expressive semantics to distinguish the current memory, during program execution, and the points where external functions are called. We call this expanded semantics *conjunctive semantics* and it extends CompCert semantics with the following

- `get_mem` and `set_mem`: The state of every language in CompCert can be interpreted as a pair of a *core* and a memory [4]. `get_mem` is the projection that returns the memory inside the state and `set_mem` changes the memory.
- `entry_points`: This is a generalization of `initial_state`, as described in section 4.1.
- `at_external` : This function exposes when a program is about to call an external function and it returns the function and the arguments being passed. The instantiation for Clight is shown in table 4.9.

Our conjunctive semantics is very closely related to *interaction semantics* [4] with two main differences. First, we don't need to define `after_external` for every language. Second, states that are `at_external` can take a step in the CompCert semantics; namely,

the execution will continue by calling the external function. The CompCert semantics, in this case, represents the *thread-local* view (or *module-local*) , where external functions, other threads and modules are abstracted into oracles that execute in one step.⁴.

In fact, given a conjunctive semantics we can derive an interaction semantics, if only we define `after_external`. The step relation is constructed by removing steps from states that make external function calls, as described by `at_external`.

4.4 More revealing simulations

CompCert’s compiler specification is stated as the following semantic preservation theorem

Theorem 4.4.1 (CompCert semantic preservation) *Let S be a source program and C its compiled version. For all behaviors B that don’t go wrong, if S has behavior B , then C also has behavior B . In short:*

$$\forall B \notin \text{Wrong}. S \Downarrow B \Rightarrow C \Downarrow B \quad (4.1)$$

Here, a behavior is a trace and a termination or divergence. If a specification *spec* is a function of behavior, then it also holds that CompCert preserves specifications in the sense that:

$$S \models \text{spec} \Rightarrow C \models \text{spec} \quad (4.2)$$

Such specification fails to preserve richer notions of specification, such as the higher-order, separation logic specifications that can be proven on Clight programs by tools like [1] or [?]. Moreover, the high level specification in 4.1 is not well suited for modular reasoning to support shared memory concurrency or compositional compilation [4].

⁴In CompCert the oracle, called `external_functions_sem`, is passed as a parameter to the correctness proof and gives the semantic of external functions.




Table 4.10: Step simulation step diagrams. (a) if s_1 takes a step to s_2 with trace t and s_1 is related to some s_2 , then there s_2 can take a number of steps with trace t to a new state s'_2 related to s'_1 . (b) The new diagram exposes the memory re-ordering injections j and j' and the traces t and t' are equivalent up to injection, by `inject_trace_strong j' t t'`.

We consider that the simulations that CompCert uses to prove 4.1 are better suited for these purposes. CompCert proves a forward simulation between its source and target executions which, together with the determinism of the target language, imply 4.1. These simulations, encoded in the record `fsim_properties`, state that (1) public global variables and functions are preserved, (2) `initial_states` are preserved (4.5), (3) `final_states` are preserved and (4) execution is preserved (4.10(a)). The simulations is parametric on a *match relation*, noted as \sim , as an invariant of related states in source and target; the relation is established at initial states and preserved by the step simulation.

For all CompCert phases, the *match state* relation describes how the memory changes after compilation. In some passes, memory doesn't change at all (e.g. Csh-mgen or Linearize) and sometimes the memory is extended by increasing the size of existing memory blocks, with new values (e.g. Allocation, Tunneling). In other cases, memory is reordered, memory blocks are coalesced, and some are unmapped. CompCert expresses this reordering with *memory injections* that map memory blocks, to their new block with some offset. For example, in Cminorgen the compiler coalesces all stack-allocate local variables of a function into a single stack block. We use this

same injection to describe how traces with memory events evolve through compilation (4.10(b)).

We propose a more expressive simulation `inject_sim` that improves the CompCert simulations in the following ways:

- Exposes how the memory changes: We expose the memory injection j that describes how memory changes after compilation. For simplicity of the proofs, for compiler passes that preserve the memory or just extend it, we also define the simpler simulations `eq_sim` and `extend_sim` respectively. These simulation follow immediately from the ones already proven in CompCert. All of the simulations we define compose horizontally to `inject_sim` as shown by the composition lemmas 4.4.2, 4.4.3 and 4.4.4.

Lemma 4.4.2 *For all semantics L_1 and L_2 if `eq_sim` L_1 L_2 then `extend_sim` L_1 L_2*

Lemma 4.4.3 *For all semantics L_1, L_2, L_3 , if `extend_sim` L_1 L_2 and `inject_sim` L_2 L_3 , then `inject_sim` L_1 L_3*

Lemma 4.4.4 *For all semantics L_1, L_2, L_3 , if `inject_sim` L_1 L_2 and `inject_sim` L_2 L_3 , then `inject_sim` L_1 L_3*

- Preserves external function calls: The original CompCert simulation only preserves traces so, for example, a compiler could replace an external function call with internal code that produces the same event. In fact the compiler does exactly that with some special external calls such as `memcpy` and certain builtins. However the compiler does not do that with arbitrary external functions (of course not!), but the simulation specification does not rule it out. We add `preserves_atx` to the simulation, which says that if a source state is




Table 4.11: At external step diagram (`simulation_atx`). Exclusive for external function calls, this diagram follows the simulation diagram in 4.10, but enforces that the compiled execution takes only one step.

`at_external`, then any target state it matches is also `at_external` with the same functions and related arguments (i.e., equal up to memory injection).

- Preserves the number of steps taken by external functions: This fact was already proven in the CompCert but was hidden in the less expressive simulation. We include a new diagram (4.11), `simulation_atx` which says that if a source state, that is `at_external` takes exactly one step then the matching target state does the same (as opposed to any number of steps as in 4.10), and the two resulting states match.

We further expand the notion of `simulation_atx` at the end of this subsection.

- Can start executions with functions that take arguments and are not `main`. We replace the `initial_state` diagram with the diagram for `entry_point` as described in 4.5.

It might be surprising that we don't further change the diagram for `entry_points`. In CompComp [4], the initial core simulation must accept almost arbitrary (but injected) memories. Our techniques allows us to assume that the the context is not changing while the program compiles. Similarly we can expect that the execution of external functions changes, based how the compiler reorders the memory, but the external function will not change the order in which it allocates memory.

```

1 Definition simulation_atx_inj_stronger {index: Type} {L1 L2: semantics}
2      (match_states: index → meminj → state L1 → state L2 →  $\mathbb{P}$ ) :=
3       $\forall s1 f \text{ args},$ 
4      at_external L1 s1 = Some (f,args) →
5       $\forall t s1' i f s2, \text{Step L1 s1 t s1}' \rightarrow$ 
6          match_states i f s1 s2 →
7           $\exists f', \text{Values.inject\_incr f f}' \wedge$ 
8           $(\exists i' s2' t',$ 
9              Step L2 s2 t' s2'  $\wedge$ 
10             match_states i' f' s1' s2'  $\wedge$ 
11             inject_trace_strong f' t t')  $\wedge$ 
12             ( $\forall t', \text{inject\_trace f' t t}' \rightarrow$ 
13                  $\exists i', \exists s2',$ 
14                 Step L2 s2 t' s2'  $\wedge$ 
15                 match_states i' f' s1' s2') .

```

Table 4.12: Stronger simulation for external steps, that universally quantifies over all injected traces. Lines 8-11 describe the existentially quantified diagram as described in 4.11. Lines 12-15, in bold, describe all the other executions that may have undefined values determined. `inject_trace` is the predicate that allows undefined values to be mapped to defined ones.

As mentioned before, our definition of `simulation_atx` might be too strong for external functions that have traces with undefined values. If those were read from memory allocated by the compiling program, it is reasonable that the values become concrete as the program compiles. Fortunately, it is enough (and easier) to prove the stricter version described in 4.11. We do provide the more permissive version of the simulation as part of the external specification of the compiler. The full Coq definition is presented in 4.12 with the addition highlighted.

Full injections

Most passes in CompCert preserve the contents in memory. Even injection passes, such as Cminorgen, Stacking and Inlining, only reorder memory and coalesce blocks, but don't remove any content from memory. Only two passes currently remove contents out of memory: SimplLocals, which pulls scalar variables whose address is not

taken into temporary variables; and `Unusedglob`, which removes unused static globals. For those injection passes where memory content is preserved, we make it explicit by adding a predicate `full_injection`, that states that an injection maps all valid blocks in memory. In the remaining of this subsection, we explain the current limitations of the way CompCert specifies unmapped parts of memory. In our version of CompCert, a compiler that skips `SimplLocals` and `Unusedglob`, can expose `full_injection` and overcome those limitations. Certainly, requiring all memory to be mapped is also a strong limitation. In what remains of this chapter, we will make the problem clear and propose a solution (although the implementation is beyond the scope of this thesis). We further discuss solutions for this limitation in related work [2](#) and in our future work [6.1](#) sections.

Consider the `remember()` and `incr()` functions from [4.1](#). As we discussed before, the execution of `incr` depends on the location in memory `*buff`. We already discussed that such functions cannot satisfy the strict "correctness" requirements of CompCert and we have corrected this problem with memory events. The second problem with this simple function, however, is that it relies on the fact that the compiler does not remove `buff` from memory. CompCert does, in fact, preserve that piece of memory, since its address has escaped, but this fact is not part of the compiler's specification.

As a second example, consider shared memory concurrency. When two threads are interacting through memory, each thread needs to know that the memory it gains access to, is not unmapped and unchanged. A thread can only use the locations it has permission over (which is a superset of the locations it accesses). This approach allows us to ensure that the memory doesn't change when other threads execute. Unfortunately, if part of the memory is unmapped, we can't ensure that the threads execute correctly. This problem is surprisingly close to the `incr()` example, and many of the solutions for that problem will also solve the problem for concurrency.

In its original paper about CompCert Leroy [2] claims that "inputs given to the programs are uniquely determined by their previous outputs". That seems to suggest that functions like `incr()` would be safe, but in its implementation CompCert rather requires that "inputs given to the programs are uniquely determined by their last outputs" (i.e. the arguments to the external function call). To create such a specification for external functions, we would create a history `args_hist` that records every argument passed to external functions. Then, external functions can depend on the entire `args_hist`. Moreover, one should be able to prove that `args_hist` are not unmapped by `SimplLocals` or `Unusedglob`, since it only contains escaping pointers. These changes are beyond the scope of the thesis, so we temporarily use `full_injection` and we skip the two problematic passes. We discuss this solution further in the 6.1 section.

Chapter 5

Compiler Correctness

This chapter describes the compiler correctness proof and its implementation.

5.1 Compiler Theorem

Here goes the compiler theorem

5.1.1 Self simulations

5.2 Compiler Theorem Implementation

Here goes the compiler theorem implementation

Chapter 6

Conclusion

Conclusion text.

6.1 Future Work

Lots to talk about here.

Appendix A

Implementation Details

Implementation details.

A.1 CPM design

Bibliography

- [1] Qinxiang Cao, Lennart Beringer, Samuel Gruetter, Josiah Dodds, and Andrew W. Appel. VST-Floyd: A separation logic tool to verify correctness of C programs. *J. Autom. Reasoning*, 61(1-4):367–422, 2018.
- [2] Xavier Leroy. Formal verification of a realistic compiler. *Communications of the ACM*, 52(7):107–115, 2009.
- [3] Xavier Leroy. The CompCert verified compiler, software and ann. proof, March 2019.
- [4] Gordon Stewart, Lennart Beringer, Santiago Cuellar, and Andrew W Appel. Compositional compcert. In *POPL*, volume 50, pages 275–287. ACM, 2015.