

CONCURRENT PERMISSION MACHINE FOR
MODULAR PROOFS OF OPTIMIZING COMPILERS
WITH SHARED MEMORY CONCURRENCY.

SANTIAGO CUELLAR

A DISSERTATION

PRESENTED TO THE FACULTY

OF PRINCETON UNIVERSITY

IN CANDIDACY FOR THE DEGREE

OF DOCTOR OF PHILOSOPHY

RECOMMENDED FOR ACCEPTANCE

BY THE DEPARTMENT OF

COMPUTER SCIENCE

ADVISER: ANDREW APPEL

SUBM. DATE

© Copyright by Santiago Cuellar, ?? copyright.

All rights reserved.

Abstract

Optimizing compilers change a program based on a formal analysis of its code and modern processors further rearrange the program order. It is hard to reason about such transformations, which makes them a source of bugs, particularly for concurrent shared-memory programs where the order of execution is critical. On the other hand, programmers should reason about their program in the source language, which abstracts such low level details.

We present the Concurrent Permission Machine (CPM), a semantic model for shared-memory concurrent programs, which is: (1) sound for high-order Concurrent Separation Logic, (2) convenient to reason about compiler correctness, and (3) useful for proving reduction theorems on weak memory models. The key feature of the CPM is that it exploits the fact that correct shared-memory programs are permission coherent: threads have (at any given time) noncompeting permission to access memory, and their load/store operations respect those permissions.

Compilers are often written with sequential code in mind, and proving the correctness of those compilers is hard enough without concurrency. Indeed, the machine-checked proof of correctness for the CompCert C compiler was a major advance in the field. Using the CPM to conveniently distinguish sequential execution from concurrent interactions, I show how to reuse the (sequential) CompCert proof, without major changes, to guarantee a stronger concurrent-permission-aware notion of correctness.

Acknowledgements

Acknowledgements...

Dedication.

Contents

Abstract	iii
Acknowledgements	iv
List of Tables	viii
List of Figures	x
1 Introduction	1
2 Related Work	2
2.1 Past work 1	2
3 Top to Bottom structure	3
3.1 Main theorem	3
3.2 The concurrent permission machine	4
3.3 Limitations of the CPM model	4
4 Making Compcert a <i>conjunctive</i> compiler	5
4.1 Passing arguments to main.	9
4.1.1 Simulating initialization preprocessing.	10
4.2 Memory events.	13
4.3 Conjunctive Semantics	15
4.4 More revealing simulations	16

5	Compiler Correctness	22
5.1	Compiler Theorem	22
5.1.1	Self simulations	22
5.2	Compiler Theorem Implementation	22
6	Conclusion	23
6.1	Future Work	23
A	Implementation Details	24
A.1	CPM design	24
	Bibliography	25

List of Tables

4.1	The function remember records the address of some buffer, and incr increments it by one.	5
4.2	The initial_state in C and Clight is a call to main . It also enforces that it takes no arguments (Tnil) and returns an integer (type_int32s). . . .	6
4.3	The events in CompCert	7
4.4	Percentage change to CompCert: changes are calculated from the number of lines added as given by running git diff between our version of CompCert and the master branch. For each feature, an estimated percent is provided.	8
4.5	The entry_point predicate in Clight	11
4.6	Part of the entry_point predicate in Mach	13
4.7	The new events in CompCert: mem_effect reflects changes to memory and delta_perm_map represents trasfer of Cur permissions.	14
4.8	at_external definition for Clight. The function checks that (1) the current state is about to make a function call, (2) that the function is an External function and, (3) that the external function cannot be inlined (The compiler is allowed to inline specific functions such as memcpy and certain builtins).	16
4.9	Step simulation step diagram	17
4.10	The new simulation step diagram	17

4.11	<code>fsim_properties_inj_relaxed</code> is identical to <code>fsim_properties_inj</code> except it uses <code>simulation_atxX</code> instead of <code>simulation_atx</code>	19
4.12	This predicate, labeled <code>ec_mem_inject</code> , is one of the properties required by <code>CompCert</code>	21

List of Figures

Chapter 1

Introduction

The main ultimate goal of this thesis is to create a modular way to prove compilers safe, even in the presence of concurrency. There are already several compilers proven safe, with machine checked proofs, that give guarantees for sequential programs. Instead of reinventing the wheel, I intend to leverage those proofs and use it to give guarantees for concurrent programs.

On top of "just a compiler proof", our system is intended to be compatible with a realistic separation logic and usable on realistic machines like X86 TSO.

The main result:

Contributions: In this thesis I have....

Chapter 2

Related Work

Here I talk about related work:

2.1 Past work 1

Topic 1.

Chapter 3

Top to Bottom structure

Note.

3.1 Main theorem

Explaining the main theorem:

Assumptions:

Theorem top2bottom_correct:

CSL_correct C_program \rightarrow

CompCert_compiler C_program = Some Asm_program \rightarrow

\forall (src_m:Memory.mem) (src_cpm:Clight.state),

CSL_init_setup C_program src_m src_cpm \rightarrow

Clight.entry_point (Clight.globalenv C_program) src_m src_cpm (main_ptr C_program) nil \rightarrow

\forall (limited_builtins:Asm_core.safe_genv x86_context.X86Context.the_ge),

asm_prog_well_formed Asm_program limited_builtins \rightarrow

\forall (U:schedule), \exists (tgt_m0 tgt_m: mem) (tgt_cpm:ThreadPool.t),

permissionless_init_machine

Asm_program limited_builtins

tgt_m0 tgt_cpm tgt_m (main_ptr C_program) nil \wedge

spinlock_safe U tgt_cpm tgt_m.

- `CSL_correct C_program`: States that the program has been proven correct in CSL, as described in [section CSL ref]
- `CompCert_compiler`: States that the CompCert compiler translates `C_program` into the assembly program `Asm_program`.
- `CSL_init_setup`: States that `src_m` and `src_cpm` are the initial memory and state as defined by the program `C_program`... According to CSL
- `Clight.entry_point` : Same as above... According to Clight.
- `limited_builtins`: Statically checkable property stating that `Asm_program` doesn't have unsupported builtins.
- `asm_prog_well_formed`: Another statically checkable property stating that the initial memory and global environment created by `Asm_program` are well formed.

Conclusions:

- `permissionless_init_machine`: There exists initial memories and state for `Asm_program`
- `spinlock_safe`: The initial state is safe and "spinlock well synchronized" as defined in [reference to spinlock]

3.2 The concurrent permission machine

Describe CPM.

3.3 Limitations of the CPM model

What are the limitations? - That coin example from Ernie Cohen. - Angel transferring too many permissions proof.

Chapter 4

Making CompCert a *conjunctive* compiler

The simple code in 4.1 communicates with its environment in two main ways: (1) it takes an address as input and (2) reads from and writes to this location to increment the value stored there. We will see how the specifications of CompCert prevents us from reasoning about such program as compilation units and as external functions and we will present solutions to these limitations.

First, CompCert can't give any guarantees about compiling the code in 4.1 because it is not a complete program. It is reasonable to expect that the function `remember` runs safely, given some assumptions (e.g. `*p` is a valid address in memory). Unfortu-

```
1 int *buff;
2 void remember(int *p){
3     buff = p;
4 }
5 void incr(void){
6     ++(*buff);
7 }
```

Table 4.1: The function `remember` records the address of some buffer, and `incr` increments it by one.

Inductive $\text{initial_state} (p: \text{program}): \text{state} \rightarrow \mathbb{P} :=$
 $| \text{initial_state_intro}: \forall b f m0,$
 $\quad \text{let } ge := \text{Genv.globalenv } p \text{ in}$
 $\quad \text{Genv.init_mem } p = \text{Some } m0 \rightarrow$
 $\quad \text{Genv.find_symbol } ge \text{ } p.(prog_main) = \text{Some } b \rightarrow$
 $\quad \text{Genv.find_funct_ptr } ge \text{ } b = \text{Some } f \rightarrow$
 $\quad \text{type_of_fundef } f = \text{Tfunction Tnil type_int32s cc_default} \rightarrow$
 $\quad \text{initial_state } p (\text{Callstate } f \text{ nil Kstop } m0).$

Table 4.2: The `initial_state` in C and Clight is a call to `main`. It also enforces that it takes no arguments (`Tnil`) and returns an integer (`type_int32s`).

nately, CompCert’s semantics assumes that a program starts executing with a call to `main()` with no arguments. In fact, the only possible initial states, are characterized by a predicate `initial_state: state \rightarrow \mathbb{P}` that takes no additional arguments. You can see an instantiation of the predicate for Clight in 4.2. So, even though CompCert correctly compiles the code, it’s specification gives no guarantees of any execution other than the one that starts by calling `main` with no arguments.

Second, CompCert’s can’t reason about programs that call `remember` and `incr` because their behavior depends on their internal state (the `buff` pointer). CompCert’s semantics allows calls to external calls that are assumed to be correct but unfortunately the specification of ”correctness” is too strict, it assumes that the function’s behavior is fully determined by (1) the state of memory, (2) the function arguments and (3) the events produced by the function. The execution of `incr` also depends on arguments passed to a previous external call (namely, `*p` passed to `remember`). One could make this explicit in the event trace but CompCert events, shown in 4.3, can only contain integers, floats or pointers to global variables.

Moreover, in the CompCert semantics, the entire behavior of external functions is bundled into one big step. From looking at an execution internal steps and external function calls are uniform. This consistency is very useful when reasoning about the compilation of the program, where we want to abstract external calls. Nevertheless,

Inductive event: Type :=	Inductive eventval: Type :=
Event_syscall: string → list eventval → eventval → event	EVint: int → eventval
Event_vload: memory_chunk → ident → ptrofs → eventval → event	EVlong: int64 → eventval
Event_vstore: memory_chunk → ident → ptrofs → eventval → event	EVfloat: float → eventval
Event_annot: string → list eventval → event	EVsingle: float32 → eventval
	EVptr_global: ident → ptrofs → eventval.

Table 4.3: The events in CompCert

when reasoning about a program in a context, it is more useful to replace the big step external calls, with their small step semantics. Regrettably, the specification of CompCert does not even guarantee that the source and target programs call the same external functions. In theory, CompCert could replace an external function call with internal steps as long as they had the same (possibly empty) trace. In practice, obviously, CompCert does not do that, but it is not exposed in its specification.

Finally, the correctness of CompCert is stated as semantic preservation theorem, where the traces are the preserved behavior. The proof of this semantic preservation, uses a forward simulation¹. We find it more convenient to expose this simulation as the specification of the compiler, deriving semantic preservation as a corollary. This way we can easily express the relation between external function calls in source and target, as simulation diagrams. Moreover, we can expose the relation between memories, in source and target.

We move, then, to lift this limitations according to the following richer notion of specification

Definition 1 (Conjunctive compiler specification) *We say that a compiler's specification is conjunctive if it satisfies the following*

¹Forward simulation and determinism of the target language implies bisimulation and thus preservation of behavior.

	Percent change
Arguments in <code>main</code>	0
Injectable Traces	0
Semantics	0
Simulations	0
Total	x

Table 4.4: Percentage change to CompCert: changes are calculated from the number of lines added as given by running `git diff` between our version of CompCert and the master branch. For each feature, an estimated percent is provided.

- *The execution of a program can start in any of its public functions and all the public functions, including `main`, can take arguments.*
- *The execution trace supports events that can describe changes to memory. Since memory may be rearranged through compilation, we call these memory events.*
- *There are functions `at_external` and `after_external` that describe states before and after external calls. `at_external` extracts the external function being called and its arguments. `after_external` constructs the state after the external function executes. Also, source and target languages operate over the same memory and the memory can be separated by `get_mem` function. We call these exposed semantics.*
- *The correctness of the compiler is stated as a forward simulation and the semantic preservation derive as a corollary. The forward simulations preserve external function calls and exposes the relation between source and target memories, throughout the execution. We call this an exposed simulation.*

In the rest of the chapter, we describe how we have improved the specification of CompCert to make it conjunctive. We first describe how to generalize `initial_state` to accept functions with arguments and other than `main`. Second we describe how to add memory events to CompCert. Then we show how to easily define exposed

semantics for every language in CompCert and finally we show how to define and prove exposed simulations.

The changes described here represent only a x% change to CompCert, as measured by running `git diff` in CompCert before and after our changes. The amount changed for every feature proposed is described in Table 4.4.

4.1 Passing arguments to main.

CompCert can compile programs where `main` takes arguments, but its correctness theorem gives no guarantees about their translation. That is because its semantic model assumes that `main` takes no arguments (See 4.2), but real C programs can take up to two arguments `argc` (argument count) and `argv` (argument vector). Also, all execution in the semantics of CompCert start with a call to `main()`, but `main` is nothing but an agreed upon term for startup code. Furthermore, the semantics assumes that the initial memory contains only the global environments. In this section we define a new predicate `entry_point`, generalizing `initial_state`, which accepts any public function, including those with arguments. The predicate also admits other memories containing, for example, the stack of other processes.

Passing arguments to main is of general interest and particularly important for our work with concurrency. Spawning new threads behaves very similar to starting a program by calling `main`: The library function that spawns a new thread `f` (e.g. `pthread_create`) must create a new stack, push the arguments to stack and then call `f` just like a program startup (which calls `_start()` and `__libc_start_main()`). Our new predicate `entry_point` is general enough to capture this two types of preprocessing. Notice that if we restricted our semantics to spawning threads with no arguments, threads wouldn't be able to share pointers and thus they would all execute in disjoint

pieces of memory with no communication. That would be a much easier and less interesting result.

Now that we defined a new starting point for executions, we must prove that compilation preserves the predicate `entry_point` in the same way that CompCert’s simulation preserves `initial_state`. The proof largely follows the simulation of internal function calls which is already proven in CompCert, so we omit the details here.

4.1.1 Simulating initialization preprocessing.

Before `main` ever starts executing, there is an initialization process that sets up the stack and the registers, before `main` is executed. The preprocessing function(s) often called `_start`, `premain` or `startup`, will create a stack, set up the arguments in the stack or in registers, set up the environment (`envp`), set up the return address (a call to `exit()`), and other bookkeeping.

One of the difficulty in simulating the initialization is that, in architectures such as *x86* in 32-bit mode, all arguments are passed in memory. As the comments in the CompCert code would put it ”Snif!”[\[1\]](#). Even architectures that allow argument passing in registers, such as *x86* in 64-bit mode, have a limited number of registers and will pass arguments on the stack after those run out. This ”pre-stack”, that contains the arguments, does not correspond to any function in the call stack of the program; it corresponds to the stack frame of `_start()`, which is part of the linked program.

Our generic predicate `entry_points: mem \rightarrow state \rightarrow val \rightarrow list val $\rightarrow \mathbb{P}$` takes a memory `m`, a state `s`, a pointer to the entry function `fun_ptr` of type `val`, and a list of arguments `args`. This predicate is language dependent, but it generally has three parts:

```

1 Inductive entry_point (ge:genv): mem → state → val → list val →  $\mathbb{P}$  :=
2 | init_core:  $\forall$  f fb m0 args targ,
3   let sg:= signature_of_type targ type_int32s cc_default in
4   type_of_fundef (Internal f) = Tfunction targ type_int32s cc_default →
5   Genv.find_funct_ptr ge fb = Some (Internal f) →
6   globals_not_fresh ge m0 →
7   Mem.mem_wd m0 →
8   Val.has_type_list args (typlist_of_typelist targ) →
9   vars_have_type (fn_vars f) targ →
10  vals_have_type args targ →
11  Mem.arg_well_formed args m0 →
12  bounded_args sg →
13  entry_point ge m0 (Callstate (Internal f) args (Kstop targ) m0).

```

Table 4.5: The `entry_point` predicate in Clight

1. Checks memory is well formed. That is, it contains no ill-formed pointers to invalid memory. CompCert generally maintains that well formed programs don't create dangling pointers pointers.²
2. Arguments are well formed. Among other things, they have the right types for the function being called, they have no ill-formed pointers, and fit in the stack.
3. Global environment is allocated correctly. This makes sure that the environment is in memory and that the function is declared as a public in the environment.

In the rest of this section, we explore the definition of `entry_points` for different languages and show how we prove that different CompCert phases preserve the predicate.

C frontend

All of the C-like languages (*Clight*, *Csharp*, *Csharpminor*) have similar `entry_point`, so we present here the one for Clight in 4.5. Lines 4-6 ensure that the environment is allocated in memory and it contains the function `f` with the right type signature. Line 7 states that the initial memory has no dangling pointers. Lines 8-11 say that

²A well formed program, should not read or write to invalid pointers. Hence, dangling pointers behave semantically as undefined values and could be modeled that way.

the arguments have the right type and have no dangling pointers. The predicate `bounded-args`, enforces that the arguments fit in the stack, which is architecture dependent. It is reasonable to replace this with a small enough bound that fits all architectures, like 4, but we keep it general here. Finally, the entry state, in line 13, is defined as a call to `f` with an empty continuation.

Our empty continuation `Kstop` takes `targs` as an argument. That is because continuations also represent the call stack; `Kstop targs` represents the "pre-stack" of `_start()` that might contain some arguments for `f`.

Register transfer languages

In the CminorGen phase, CompCert coalesces all function variables into a stackframe. Some functions might get empty stackframes (i.e., a memory block with empty permissions, that cannot be written to), if none of their variables has their address taken. These stackframes are important, even the empty ones, because that is where spill variables will be written after register allocation in the Allocation phase. We follow suit and create an empty stackframe for `_start()`. The stack is empty because the compiler has not yet decided what arguments will be passed in memory and which ones in register. Even for architectures that pass all arguments in memory, this is not done until the Stacking pass. So for languages before Stacking (Cminor, CminorSel, RTL, LTL, Linear), there is an extra line in `entry_point` to make sure that the empty stackframe is allocated:

```
Mem.alloc m0 0 0 = (m1, stk)
```

The rest of the predicate is almost identical to the one in [4.5](#).

Machine languages

In the machine languages (Mach and Asm), a function expects certain shape from the stackframe of its caller. We replace the empty stackframe allocation above, with

```

1      :
2      let '(stk_sz,ret_ofs,parent_ofs) := stack_defs (fn_sig f) in
3      Mem.alloc m0 0 stk_sz = (m1, spb) →
4      let sp:= Vptr spb Ptrofs.zero in
5      store_stack m1 sp Tptr parent_ofs Vnullptr = Some m2 →
6      store_stack m2 sp Tptr ret_ofs Vnullptr = Some m3 →
7      make_arguments (Regmap.init Vundef) m3 sp
8              (loc_arguments (funsig (Internal f))) args = Some (rs, m4) →
9      :

```

Table 4.6: Part of the `entry_point` predicate in Mach

the construction of the "pre-stackframe" as shown in 4.6. The function `stack_defs` is an architecture dependent function that calculates the layout of the "pre-stack" and returns the size `stk_sz`, the offset of the return address `ret_ofs` and a back link to parent frame `parent_ofs`. The last two values are unused, but the stack must have space for them. Line 3 allocates the stack of the correct size. Lines 5-8 store the return address, the link to the parent and the arguments in the stack.

4.2 Memory events.

The visible behavior of the CompCert semantics (for all languages) is a trace of events, as described in 4.3. It records interactions with the outside world; for example, the results of a read system call will record a `Event_syscall` together with the name of the system call, its parameters, and its result. The only pointers that can appear in the trace are locations of global variable. As we explained above, these events are insufficient to capture the behavior of `incr` in 4.1. The events are also not sufficient to express any kind of behavior that depends on the memory. For example, a call to `pthread_mutex_unlock(&l)` not only changes the the state of the lock `l`, conceptually it gives away control to the data in memory protected by `l`. Such behavior that reveals locations in memory cannot be expressed in CompCert events that can't have

Inductive event: Type :=	Inductive mem_effect :=
...	Write : $\forall (b : \text{block}) (ofs : Z)$
Event_acq_rel: list mem_effect \rightarrow	(bytes : list memval), mem_effect
delta_perm_map \rightarrow list mem_effect \rightarrow	Alloc: $\forall (b : \text{block})(lo\ hi:Z), \text{mem_effect}$
event	Free: $\forall (l: \text{list } (\text{block} * Z * Z)), \text{mem_effect}.$
Event_spawn: block \rightarrow delta_perm_map \rightarrow	
delta_perm_map \rightarrow event.	

Table 4.7: The new events in CompCert: **mem_effect** reflects changes to memory and **delta_perm_map** represents transfer of **Cur** permissions.

pointers on them (except the location of global variables). More generally, these events are poorly suited to express any sort of shared memory interactions, such as concurrency or separated compilation. We propose to include 2 new types of events which we call *memory events* as described in table 4.7. As their name suggests, they contain references to locations in memory, beyond the global variables. The first one, **Event_acq_rel**, represents a generic memory interactions where the external function performs some arbitrary changes to memory, recorded by a list of **mem_effects**, and transfers some permissions recorded by **delta_perm_map**. We find it convenient to split the effects on memory as those that happen before the changes in permissions and those that happen after. The second one, **Event_spawn**, represents the creation of a new thread or a new module. It records the function being called, as a block number, and the change in permissions by two **delta_perm_map**, one representing the permissions given and the other the starting permissions of a new thread/module.

The correctness of the CompCert compiler is formulated as a preservation of traces. However, if our traces contain memory events and the memory can be reordered by compilation, the trace must be reordered accordingly. Thus our new correctness will be formulated up to memory reordering as shown in 4.9. In fact, if a compiler pass (or passes) don't change the order of memory (such as equality and extension passes), then the trace is preserved and nothing changes in the simulation. If the pass changes the order of memory according to some j , then we need show that the source and target traces, t and t' , are related by **inject_trace_strong** $j\ t\ t'$. That means that events

in t and t' are identical except all memory locations are reordered according to j . To add this property, we need simulations to expose how they reorder memory. The full description of our new conjunctive simulation will be explained in section 4.4.

Compilers not only reorder memory, sometimes undefined values in the source program are mapped to concrete values in the target. Such mappings are useful in compiler passes such as register coalescing, where registers that were previously uninitialized, can now map to an initialized register with concrete values. The predicate `inject_trace_strong` always maps undefined values to undefined values. This is a reasonable restriction since inspecting undefined values is not an allowed behavior so they shouldn't appear in the trace. However, if an application required traces with undefined values, we can still support that. It turns out that it is enough to prove that the execution with the *strongly injected* trace is safe, and from it we can derive safe executions for all traces that have some undefined values defined.

4.3 Conjunctive Semantics

As described in the introduction, we need more expressive semantics to distinguish the current memory, during program execution, and the points where external functions are called. We call this expanded semantics *conjunctive semantics* and it extends CompCert semantics with the following

- `get_mem` and `set_mem`: The state of every language in CompCert can be interpreted as a pair of a *core* and a memory [2]. `get_mem` is the projection that returns the memory inside the state and `set_mem` changes the memory.
- `entry_points`: This is a generalization of `initial_state`, as described in section 4.1.
- `at_external` : This function exposes when a program is about to call an external function and it returns the function and the arguments being passed. The instantiation for Clight is shown in table 4.8.

Definition `at_external (c: state) : option (external_function * list val) :=`
`match c with`
`| Callstate fd args k _ =>`
`match fd with`
`| External ef targ tres cc => if ef_inline ef then None else Some (ef, args)`
`| _ => None`
`end`
`| _ => None`
`end.`

Table 4.8: `at_external` definition for Clight. The function checks that (1) the current state is about to make a function call, (2) that the function is an External function and, (3) that the external function cannot be inlined (The compiler is allowed to inline specific functions such as `memcpy` and certain builtins).

Our conjunctive semantics is very closely related to *interaction semantics* [2] with two main differences. First, we don't need to define `after_external` for every language. Second, states that are `at_external` can take a step in the CompCert semantics; namely, the execution will continue by calling the external function. The CompCert semantics, in this case, represents the *thread-local* view (or *module-local*), where external functions, other threads and modules are abstracted into oracles that execute in one step.³

In fact, given a conjunctive semantics we can derive an interaction semantics, if only we define `after_external`. The step relation is constructed by removing steps from states that make external function calls, as described by `at_external`.

4.4 More revealing simulations

We extend CompCert simulations in several ways:

- Expose the injection, use `get_mem` to show that the memories are either related by an injection, an extension or are equal. We provide three types of

³In CompCert the oracle, called `external_functions_sem`, is passed as a parameter to the correctness proof and gives the semantic of external functions.

$$\begin{aligned}
& \forall s_1 \ t \ s'_1 \ f. \ s_1 \xrightarrow{t} s'_1 \rightarrow \\
& \quad \forall i \ s_2, \ s_1 \xrightarrow{f,i} s_2 \rightarrow \\
& \quad \exists i' \ s'_2 \ f' \ t', \\
& (s_2 \xrightarrow{t'}^+ s'_2) \vee (s_2 \xrightarrow{t'}^* s'_2 / i' \leq i)) \wedge \\
& \quad s_1 \xrightarrow{f',i'} s_2 \wedge \\
& \quad f \sqsubseteq f' \wedge \\
& \text{inject_trace_strong } f' \ t \ t'
\end{aligned}$$

Table 4.9: Step simulation step diagram

`fsim.simulation`:

$$\begin{aligned}
& \forall s_1 \ t \ s'_1 \ f, \text{Step L1 } s_1 \ t \ s'_1 \rightarrow \\
& \quad \forall i \ s_2, \text{match_states } i \ f \ s_1 \ s_2 \rightarrow \\
& \quad \exists i', \exists s'_2 \ f' \ t', \\
& (\text{Plus L2 } s_2 \ t' \ s'_2 \vee (\text{Star L2 } s_2 \ t' \ s'_2 \wedge i' \leq i)) \\
& \text{match_states } i' \ f' \ s'_1 \ s'_2 \wedge \\
& \text{inject_incr } f \ f' \wedge \\
& \text{inject_trace_strong } f' \ t \ t'
\end{aligned}$$

Table 4.10: The new simulation step diagram

simulations for injections, extensions and equalities named `eq_sim`, `extend_sim` and `inject_sim` respectively. They all compose to injections as shown by the composition lemmas 4.4.1, 4.4.2 and 4.4.3:

Lemma 4.4.1 *For all semantics L_1 and L_2 if `eq_sim` $L_1 \ L_2$ then `extend_sim` $L_1 \ L_2$*

Lemma 4.4.2 *For all semantics L_1, L_2, L_3 , if `extend_sim` $L_1 \ L_2$ and `inject_sim` $L_2 \ L_3$, then `inject_sim` $L_1 \ L_3$*

Lemma 4.4.3 *For all semantics L_1, L_2, L_3 , if `inject_sim` $L_1 \ L_2$ and `inject_sim` $L_2 \ L_3$, then `inject_sim` $L_1 \ L_3$*

- We also need to know that the compiler doesn't change the number of steps taken by external functions. This fact was already proven in the CompCert

correctness proof but was hidden by a less expressive simulation, so we add it to the simulation as `simulation_atx`. The new fact says that if a source state, that is `at_external`, takes exactly one step then the matching target state does the same (as opposed to any number of steps), and the two resulting states match. From the match relation, we can also derive that if the resulting source state is in `after_external`, then so is the target one.

This fact is crucial to go from thread-local simulation to full-program simulation, where we need to replace the "oracular" external steps by their real execution of the external function.

We further expand the notion of `simulation_atx` at the end of this subsection.

- We need to know that the compiler preserves external function calls. The original CompCert simulation only preserves traces so, for example, a compiler could replace an external function call with internal code that produces the same event. This is not what the compiler does, but the simulation proof does not rule it out. We add `preserves_atx` to the simulation, which says that if a source state is `at_external`, then any target state it matches is also at external.

It might be surprising that we don't change the the simulation diagram for `initial_state` (other than changing it to `entry_points`). In CompComp [2], the initial core simulation must accept almost arbitrary (but injected) memories. Our technique allows us to assume that the the context is not compiled, and thus the initial memory is identical in source and target.

Some readers might be surprise by the shape of `simulation_atx`; a simulation for external functions should be universally quantified, instead of existentially. That is, the simulation should accept any trace of the external world (up to injection) and not just one. CompCert can get away with this form because traces are identical in source and target, so only one trace needs to be accounted for. In our setting, traces

Definition `simulation_atx_inj_relaxed` {index:Type} {L1 L2: semantics}
 $(\text{match_states: index} \rightarrow \text{meminj} \rightarrow \text{state L1} \rightarrow \text{state L2} \rightarrow \mathbb{P}) :=$
 $\forall s1 \ f \ \text{args},$
 $\text{at_external L1 } s1 = \text{Some } (f, \text{args}) \rightarrow$
 $\forall t \ s1' \ i \ f \ s2, \text{Step L1 } s1 \ t \ s1' \rightarrow$
 $\text{match_states } i \ f \ s1 \ s2 \rightarrow$
 $\exists f', \text{Values.inject_incr } f \ f' \wedge$
 $(\exists i' \ s2' \ t',$
 $\text{Step L2 } s2 \ t' \ s2' \wedge$
 $\text{match_states } i' \ f' \ s1' \ s2' \wedge$
 $\text{inject_trace_strong } f' \ t \ t') \wedge$
 $\forall t', \text{inject_trace } f' \ t \ t' \rightarrow$
 $\exists i', \exists s2',$
 $\text{Step L2 } s2 \ t' \ s2' \wedge$
 $\text{match_states } i' \ f' \ s1' \ s2'.$

Stronger simulation for external steps, that universally quantifies over all injected traces.

Lemma `injection_injection_relaxed_composition`:

$\forall (L1 \ L2 \ L3 : \text{semantics}),$
 $\text{fsim_properties_inj } L1 \ L2 \rightarrow$
 $\text{fsim_properties_inj_relaxed } L2 \ L3 \rightarrow$
 $\text{fsim_properties_inj_relaxed } L1 \ L3.$

Table 4.11: `fsim_properties_inj_relaxed` is identical to `fsim_properties_inj` except it uses `simulation_atxX` instead of `simulation_atx`.

are injected (`inject $t_s \ t_t$`), and this relation is not 1-to-1 because an injection maps undefined values to any value. So, if we expect to use the compiler in a real context, we need to expand the definition of `simulation_atx` to universally quantify over traces, as done in figure 4.4. Does this mean we need to change the proofs for all injection phases? No! The relations `simulation_atx` and `simulation_atxX` compose horizontally into `simulation_atxX`, as described in the lemma `injection_injection_relaxed_composition`. This means that we can prove CompCert with respect to `simulation_atx` and use the selfsimulation of Assembly (5.1.1) to prove the stronger simulation. The composed simulation will result in `simulation_atxX` as desired.

Full injections

Most passes in CompCert preserve the contents in memory. Even injection passes, such as Cminorgen, Stacking and Inlining, only reorder memory and coalesce blocks, but don't remove any content from memory. Only two passes currently pull contents out of memory: SimplLocals, which pulls scalar variables whose address is not into temporary variables, and Unusedglob, which removes unused static globals. For those injection passes where memory content is preserved, we make that explicit by adding a predicate `full_injection`, that states that an injection maps all valid blocks in memory. In the remaining of this subsection, we explain the current limitations of the way CompCert specifies unmapped parts of memory. In our version of CompCert, a compiler that skips SimplLocals and Unusedglob, can expose `full_injection` and overcome those limitations. Certainly, requiring all memory to be mapped is also a strong limitation. We discuss solutions for this limitation in related work [2](#) and in our future work [6.1](#).

Consider the following possible system call that executes a `write` in two steps: a first call to `twostep_write_set_buffer(int fd, const void *buf, size_t nbytes)` records the arguments and a second call to `twostep_write()` executes the write, which is equivalent to calling the `write` system call with the arguments recorded in the first call.

Notice that the safe execution of `twostep_write()` depends on the buffer `*buf` being still available in memory. Unfortunately, CompCert imposes requirements on external functions that rule this out. In particular, CompCert requires the external function to commute with memory injections (as described in [4.4](#)) with the only knowledge that memories are injected, arguments are injected and symbols are preserved. These conditions do not ensure that `*buf` is in memory. Since `*buf` is not in the arguments, it could be unmapped. Knowing that the injection is `full_injection`, is enough to know that `*buf` is not unmapped.

$$\begin{aligned}
& \forall \text{ge1 ge2 vars m1 t vres m2 f m1' vars'}, \\
& \text{symbols.inject f ge1 ge2} \rightarrow \\
& \text{sem ge1 vars m1 t vres m2} \rightarrow \\
& \text{Mem.inject f m1 m1'} \rightarrow \\
& \text{Val.inject_list f vars vars'} \rightarrow \\
& \exists f', \exists \text{vres'}, \exists \text{m2'}, \exists t', \\
& \quad \text{sem ge2 vars' m1' t' vres' m2'} \\
& \quad \wedge \text{Val.inject f' vres vres'} \\
& \quad \wedge \text{Mem.inject f' m2 m2'} \\
& \quad \wedge \text{Mem.unchanged_on (loc.unmapped f) m1 m2} \\
& \quad \wedge \text{Mem.unchanged_on (loc.out_of_reach f m1) m1' m2'} \\
& \quad \wedge \text{inject_incr f f'} \\
& \quad \wedge \text{inject_separated f f' m1 m1'};
\end{aligned}$$

Table 4.12: This predicate, labeled `ec_mem.inject`, is one of the properties required by CompCert.

As a second example, consider shared memory concurrency. When two threads are interacting through memory, each thread needs to know that the memory it gains access to, is not unmapped and unchanged. We specify the memory used by a thread by the locations it has permissions over (which is a superset of the locations it access). This approach allows us to ensure that the memory doesn't change when other threads execute. Unfortunately, if part of the memory is unmapped, we can't ensure that the threads execute correctly. This problem is surprisingly close to the `twostep-write()` example, and many of the solutions for that problem will also solve the problem for concurrency. Temporarily, having `full_injection` allows for concurrency.

Chapter 5

Compiler Correctness

This chapter describes the compiler correctness proof and its implementation.

5.1 Compiler Theorem

Here goes the compiler theorem

5.1.1 Self simulations

5.2 Compiler Theorem Implementation

Here goes the compiler theorem implementation

Chapter 6

Conclusion

Conclusion text.

6.1 Future Work

Lots to talk about here.

Appendix A

Implementation Details

Implementation details.

A.1 CPM design

Bibliography

- [1] Xavier Leroy. The CompCert verified compiler, software and ann. proof, March 2019.
- [2] Gordon Stewart, Lennart Beringer, Santiago Cuellar, and Andrew W Appel. Compositional compcert. In *POPL*, volume 50, pages 275–287. ACM, 2015.